

Anthony Mendez | 862307065 | amend303

Jordan Kushner | 862294132 | jkusc002

Due Date: 06/08/23

Lab 3 - XV6 Threads

Video Demo Link: <https://youtu.be/2ZpMW1UNk6g>

| | |
|---|-----------|
| All Modified Files..... | 2 |
| Change Explanations and Screenshots..... | 3 |
| Description of XV6 Source Code..... | 13 |
| Summary of Contributions..... | 14 |

All Modified Files

- Kernel/ files
 - Syscall.h
 - Syscall.c
 - Sysproc.c
 - Proc.c
 - Defs.h
 - Proc.h
 - Trampoline.s
 - Trap.c
- User/ files
 - Usys.pl
 - User.h
 - Lab3_test.c
 - Thread.h
 - Thread.c
- Makefile

Change Explanations and Screenshots

- Syscall.h

- Syscall number

- ```
23 #define SYS_clone 22 //edited
```

- Syscall.c

- System call function prototype

- ```
104  extern uint64 sys_clone(void); //edited
```

- Syscall number to array mapping

- ```
130 [SYS_clone] sys_clone, //edited
```

- Sysproc.c

- Helper function to call syscall

- ```
30  uint64
31  sys_clone(void) //edited
32  {
33      uint64 addr;
34      addr = myproc()->kstack;
35      //argaddr(0, &addr);
36      return clone((void*)addr);
37  }
```

- Proc.c

- Clone system call to create a new thread in a process. Similar to fork().

Calls `allocproc_thread` to create a child thread, sets the trapframe to the stack argument, shares the parents pagetable, maps the trapframe, and sets the state as runnable. Returns child pid to the calling process and returns 0 to child thread.

```

392 int //edited
393 clone(void* stack)
394 {
395     int i, pid;
396     struct proc *np;
397     struct proc *p = myproc();
398
399     // Each process has at most 20 threads
400     if (p->threadCounter > 20) {
401         return -1;
402     }
403
404     // Allocate thread
405     if((np = allocproc_thread()) == 0){
406         return -1;
407     }
408
409     if(stack == 0){
410         return -1;
411     }
412
413     np->trapframe->sp = (uint64)stack;
414     np->pagetable = p->pagetable;
415     np->sz = p->sz;
416
417     if(mappages(np->pagetable, TRAPFRAME - (PGSIZE * np->thread_id), PGSIZE,
418         (uint64)(np->trapframe), PTE_R | PTE_W) < 0){
419         uvmunmap(np->pagetable, TRAMPOLINE, 1, 0);
420         uvmfree(np->pagetable, 0);
421         return -1;
422     }
423
424     // copy saved user registers.
425     *(np->trapframe) = *(p->trapframe);
426
427     // Cause fork to return 0 in the child.
428     np->trapframe->a0 = 0;

```

```

428 // increment reference counts on open file descriptors.
429 for(i = 0; i < NOFILE; i++)
430 |   if(p->ofile[i])
431 |     np->ofile[i] = filedup(p->ofile[i]);
432 np->cwd = idup(p->cwd);
433
434 safestrcpy(np->name, p->name, sizeof(p->name));
435
436 pid = np->pid;
437
438 release(&np->lock);
439
440 acquire(&wait_lock);
441 np->parent = p;
442 release(&wait_lock);
443
444 acquire(&np->lock);
445 np->state = RUNNABLE;
446 release(&np->lock);
447
448 // Return the thread id to the parent process, return 0 to the child
449 if(p->thread_id == 0) {
450 |   return pid;
451 | }
452
453 return 0;
454 }

```

○

- Allocproc_thread to allocate the page and begin running the new thread.

Mostly the same as allocproc() except it does not allocate a separate page table.

```

350 static struct proc*
351 allocproc_thread(void)
352 {
353     struct proc *p;
354
355     for(p = proc; p < &proc[NPROC]; p++) {
356         acquire(&p->lock);
357         if(p->state == UNUSED) {
358             goto found;
359         } else {
360             release(&p->lock);
361         }
362     }
363     return 0;
364
365 found:
366     p->pid = allocpid();
367     p->state = USED;
368
369     // Allocate a trapframe page.
370     if((p->trapframe = (struct trapframe *)kalloc()) == 0){
371         freeproc(p);
372         release(&p->lock);
373         return 0;
374     }
375
376     // Set up new context to start executing at forkret,
377     // which returns to user space.
378     memset(&p->context, 0, sizeof(p->context));
379     p->context.ra = (uint64)forkret;
380     p->context.sp = p->kstack + PGSIZE;
381
382     // Assign the thread id and increment the thread counter.
383     p->thread_id = threadCounter++;
384
385     return p;
386 }

```

○

- Freeproc deallocation of child resources after the wait system call.

Unmapping the thread based on its location in the stack and decreasing the counter since the thread is free'd.

```

163 static void
164 freeproc(struct proc *p)
165 {
166     if (p->trapframe) {
167         kfree((void*)p->trapframe);
168     }
169     p->trapframe = 0;
170     if (p->thread_id == 0) {
171         if(p->pagetable) {
172             proc_freepagetable(p->pagetable, p->sz);
173         }
174         p->pagetable = 0;
175     }
176     else {
177         // Free the thread trapframe page and decrement the thread counter
178         uvmunmap(p->pagetable, TRAPFRAME - (PGSIZE * p->thread_id), 1, 0);
179         p->threadCounter--;
180     }
181
182     p->sz = 0;
183     p->pid = 0;
184     p->parent = 0;
185     p->name[0] = 0;
186     p->chan = 0;
187     p->killed = 0;
188     p->xstate = 0;
189     p->state = UNUSED;
190 }

```

○

- Allocproc initializing thread count and thread id

```

150 // Set up new context to start executing at forkret,
151 // which returns to user space.
152 memset(&p->context, 0, sizeof(p->context));
153 p->context.ra = (uint64)forkret;
154 p->context.sp = p->kstack + PGSIZE;
155 p->threadCounter = 0;
156 p->thread_id = p->threadCounter++;
157
158 return p;
159 }

```

○

- Defs.h

- System call function header

- ```
88 int clone(void*); //edited
```

- Proc.h

- Thread id for each process control block and thread counter

```
94 int thread_id; //thread id for clone
95 int threadCounter; //thread counter for clone
```

- 

- Trampoline.s

- Changed original code to code provided in lab

```
11 .section trampsec
12 .globl trampoline
13 trampoline:
14 .align 4
15 .globl uservec
16 uservec:
17 #
18 # trap.c sets stvec to
19 # traps from user space
20 # in supervisor mode, to
21 # user page table.
22 #
23 # sscratch points to what
24 # mapped into user space
25 #
26 #
27 # swap a0 and sscratch
28 # so that a0 is TRAPFRAME
29 csrrw a0, sscratch, a0
30 #
31 # save the user registers
32 sd ra, 40(a0)
33 sd sp, 48(a0)
34 sd gp, 56(a0)
35 sd tp, 64(a0)
36 sd t0, 72(a0)
37 sd t1, 80(a0)
38 sd t2, 88(a0)
39 sd s0, 96(a0)
40 sd s1, 104(a0)
41 sd a1, 120(a0)
42 sd a2, 128(a0)
43 sd a3, 136(a0)
44 sd a4, 144(a0)
45 sd a5, 152(a0)
46 sd a6, 160(a0)
47 sd a7, 168(a0)
48 sd s2, 176(a0)
49 sd s3, 184(a0)
50 sd s4, 192(a0)
51 sd s5, 200(a0)
52 sd s6, 208(a0)
53 sd s7, 216(a0)
54 sd s8, 224(a0)
55 sd s9, 232(a0)
56 sd s10, 240(a0)
57 sd s11, 248(a0)
58 sd t3, 256(a0)
59 sd t4, 264(a0)
60 sd t5, 272(a0)
61 sd t6, 280(a0)
```

-



```

64 csrr t0, sscratch
65 sd t0, 112(a0)
66
67 # restore kernel stack pointer
68 ld sp, 8(a0)
69
70 # make tp hold the current thread pointer
71 ld tp, 32(a0)
72
73 # load the address of user page table
74 ld t0, 16(a0)
75
76 # restore kernel page table
77 ld t1, 0(a0)
78 csrw satp, t1
79 sfence.vma zero, zero
80
81 # a0 is no longer valid,
82 # table does not speciall
83
84 # jump to usertrap(), whi
85 jr t0
86
87 .globl userret
88 userret:
89 # userret(TRAPFRAME, page
90 # switch from kernel to u
91 # usertrapret() calls her
92 # a0: TRAPFRAME, in user
93 # a1: user page table, fo
94
95 # switch to the user page
96 csrw satp, a1
97 sfence.vma zero, zero
98
99 # put the saved user a0 i
100 # can swap it with our a0
101 ld t0, 112(a0)
102 csrw sscratch, t0

```

```

104 # restore all but a0 fr
105 ld ra, 40(a0)
106 ld sp, 48(a0)
107 ld gp, 56(a0)
108 ld tp, 64(a0)
109 ld t0, 72(a0)
110 ld t1, 80(a0)
111 ld t2, 88(a0)
112 ld s0, 96(a0)
113 ld s1, 104(a0)
114 ld a1, 120(a0)
115 ld a2, 128(a0)
116 ld a3, 136(a0)
117 ld a4, 144(a0)
118 ld a5, 152(a0)
119 ld a6, 160(a0)
120 ld a7, 168(a0)
121 ld s2, 176(a0)
122 ld s3, 184(a0)
123 ld s4, 192(a0)
124 ld s5, 200(a0)
125 ld s6, 208(a0)
126 ld s7, 216(a0)
127 ld s8, 224(a0)
128 ld s9, 232(a0)
129 ld s10, 240(a0)
130 ld s11, 248(a0)
131 ld t3, 256(a0)
132 ld t4, 264(a0)
133 ld t5, 272(a0)
134 ld t6, 280(a0)
135
136 # restore user a0, and
137 csrrw a0, sscratch, a0
138
139 # return to user mode a
140 # usertrapret() set up
141 sret

```

- Trap.c

- In usertrapret(), telling kernel trapframe locations of children

```
130 if(p->thread_id == 0) {
131 ((void (*)(uint64,uint64))trampoline_userret)(TRAPFRAME, satp);
132 }
133 else {
134 ((void (*)(uint64,uint64))trampoline_userret)
135 (TRAPFRAME - (PGSIZE * p->thread_id), satp);
136 }
137 }
```

- Usys.pl

- User system call wrapper entry

- ```
20     entry("clone"); # //edited
```

- User.h

- Clone system call in user space

- ```
5 int clone(void*); //edited
```

- Lab3\_test.c

- Provided code to test the clone system call

```

1 #include "kernel/stat.h"
2 #include "kernel/types.h"
3 #include "user/thread.h"
4 #include "user/user.h"
5 struct lock_t lock;
6 int n_threads, n_passes, cur_turn, cur_pass;
7 void *thread_fn(void *arg) {
8 int thread_id = (uint64)arg;
9 int done = 0;
10 while (!done) {
11 lock_acquire(&lock);
12 if (cur_pass >= n_passes)
13 done = 1;
14 else if (cur_turn == thread_id) {
15 cur_turn = (cur_turn + 1) % n_threads;
16 printf("Round %d: thread %d is passing the token to thread %d\n",
17 ++cur_pass, thread_id, cur_turn);
18 }
19 lock_release(&lock);
20 sleep(0);
21 }
22 return 0;
23 }
24 int main(int argc, char *argv[]) {
25 if (argc < 3) {
26 printf("Usage: %s [N_PASSES] [N_THREADS]\n", argv[0]);
27 exit(-1);
28 }
29 n_passes = atoi(argv[1]);
30 n_threads = atoi(argv[2]);
31 cur_turn = 0;
32 cur_pass = 0;
33 lock_init(&lock);
34 for (int i = 0; i < n_threads; i++) {
35 thread_create(thread_fn, (void *) (uint64)i);
36 }
37 for (int i = 0; i < n_threads; i++) {
38 wait(0);
39 }
40 printf("Frisbee simulation has finished, %d rounds played in total\n",
41 n_passes);
42 exit(0);
43 }

```

- Thread.h

- Header file with lock structure and function prototypes.

```

1 struct lock_t {
2 uint locked;
3 };
4
5
6 int thread_create(void *(start_routine)(void*), void *arg);
7 void lock_init(struct lock_t* lock);
8 void lock_acquire(struct lock_t* lock);
9 void lock_release(struct lock_t* lock);
10

```

- Thread.c
  - Lock function implementations and thread creation by passing stack. Lock init begins locked. Acquire holds the lock in place and syncs it. Release syncs and releases the lock. The return pid will be 0 for the child and the routine starts and exits for the child. The parent receives the child's pid

```

1 #include "kernel/stat.h"
2 #include "kernel/types.h"
3 #include "user/thread.h"
4 #include "user/user.h"
5
6 int thread_create(void *(start_routine)(void*), void *arg){
7 void* stack = malloc(4096);
8 int cloneVal = clone(stack);
9 if(cloneVal == -1){
10 return -1;
11 }else if(cloneVal == 0){
12 start_routine(arg);
13 exit(0);
14 }else{
15 return 0;
16 }
17 }
18
19 void lock_init(struct lock_t* lock){
20 lock->locked = 0;
21 }
22
23 void lock_acquire(struct lock_t* lock){
24 while(__sync_lock_test_and_set(&lock->locked, 1) != 0)
25 ;
26 __sync_synchronize();
27 }
28
29 void lock_release(struct lock_t* lock){
30 __sync_synchronize();
31 __sync_lock_release(&lock->locked);
32 }

```

- Makefile
  - Added lab3\_test into programs

- `135 $U/_lab3_test\`
- Thread for the thread library
- `90 ULIB = $U/ulib.o $U/usys.o $U/printf.o $U/umalloc.o $U/thread.o`

## Description of XV6 Source Code

The default XV6 code supports process creation, but it does not support threads. In this lab, we implement the clone system call to support thread creation. The clone system call allocates a thread for the calling process and returns the child pid to the parent and 0 to the child. We also implemented a user level thread library to clone child threads with user stacks. The thread library also uses a lock that can be acquired and released. The thread create function creates a stack the size of a page, passes it to the clone function, and begins the thread execution until it finishes and exits. The test begins in the user space by calling the thread create function. Then the thread create function calls the clone function and begins the thread execution. Once the clone function is called, the system call is identified in the user space and control is transferred over to the kernel space. The clone call allocates the thread, does the required mappings and proc assignments, and returns the child thread pid to the parent process.

## Summary of Contributions

Jordan implemented the thread library, trampoline, trap code, `allocproc_thread()` and `clone()`. Anthony also worked on the `clone()` system call, helped debug some important “panic: “ type errors, and wrote the report.